

titanic-dataset-my-first-submission

August 12, 2022

0.1 Titanic Dataset

My goal is build a new version ¹ of predictive model that predicts which passengers survived the Titanic shipwreck; and try to use some "advanced" (for me :smile:) techniques such as cross validation, grid search and an ensemble algorithm like "Random Forest Classifier". Let's see what happen :smile:

1 Import data set

```
[ ]: import pandas as pd

# We'll use a dataset taken from: https://www.kaggle.com/competitions/titanic
dfTrain = pd.read_csv("./data/train.csv", sep=',')
dfTest = pd.read_csv("./data/test.csv", sep=",")
```

2 Basic EDA and cleaning data

```
[ ]: from basic_exploration import *
basicEDA(dfTrain, "Titanic Train")
```

Just first five rows

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	

¹See <https://www.kaggle.com/code/francescopaolol/logisticregression-on-complete-titanic-dataset>

3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1
4	Allen, Mr. William Henry	male	35.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

'The Titanic Train data set consists of 12 different features which for 891 samples.'

Info about the index dtype and columns, non-null values and memory usage.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null    int64
1   Survived         891 non-null    int64
2   Pclass           891 non-null    int64
3   Name             891 non-null    object
4   Sex              891 non-null    object
5   Age              714 non-null    float64
6   SibSp            891 non-null    int64
7   Parch            891 non-null    int64
8   Ticket           891 non-null    object
9   Fare             891 non-null    float64
10  Cabin            204 non-null    object
11  Embarked         889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

None
```

Check all of the values in the data frame which holds my data.

	PassengerId	Survived	Pclass	Name \
0	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
1	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
2	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
3	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
4	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
..
886	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
887	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>
888	<class 'int'>	<class 'int'>	<class 'int'>	<class 'str'>

```

889 <class 'int'> <class 'int'> <class 'int'> <class 'str'>
890 <class 'int'> <class 'int'> <class 'int'> <class 'str'>

      Sex      Age      SibSp      Parch \
0  <class 'str'> <class 'float'> <class 'int'> <class 'int'>
1  <class 'str'> <class 'float'> <class 'int'> <class 'int'>
2  <class 'str'> <class 'float'> <class 'int'> <class 'int'>
3  <class 'str'> <class 'float'> <class 'int'> <class 'int'>
4  <class 'str'> <class 'float'> <class 'int'> <class 'int'>
..      ...      ...      ...      ...
886 <class 'str'> <class 'float'> <class 'int'> <class 'int'>
887 <class 'str'> <class 'float'> <class 'int'> <class 'int'>
888 <class 'str'> <class 'float'> <class 'int'> <class 'int'>
889 <class 'str'> <class 'float'> <class 'int'> <class 'int'>
890 <class 'str'> <class 'float'> <class 'int'> <class 'int'>

      Ticket      Fare      Cabin      Embarked
0  <class 'str'> <class 'float'> <class 'float'> <class 'str'>
1  <class 'str'> <class 'float'> <class 'str'> <class 'str'>
2  <class 'str'> <class 'float'> <class 'float'> <class 'str'>
3  <class 'str'> <class 'float'> <class 'str'> <class 'str'>
4  <class 'str'> <class 'float'> <class 'float'> <class 'str'>
..      ...      ...      ...      ...
886 <class 'str'> <class 'float'> <class 'float'> <class 'str'>
887 <class 'str'> <class 'float'> <class 'str'> <class 'str'>
888 <class 'str'> <class 'float'> <class 'float'> <class 'str'>
889 <class 'str'> <class 'float'> <class 'str'> <class 'str'>
890 <class 'str'> <class 'float'> <class 'float'> <class 'str'>

```

[891 rows x 12 columns]

Count na values

```

PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64

```

"The columns with missing data are: ['Age', 'Cabin', 'Embarked']"

Percent of missing 'Age' records is 19.865 %

Percent of missing 'Cabin' records is 77.104 %

Percent of missing 'Embarked' records is 0.224 %

Show the statistic report of the numeric features of the dataset

	count	mean	std	min	25%	50%	75%	\
PassengerId	891.0	446.000000	257.353842	1.00	223.5000	446.0000	668.5	
Survived	891.0	0.383838	0.486592	0.00	0.0000	0.0000	1.0	
Pclass	891.0	2.308642	0.836071	1.00	2.0000	3.0000	3.0	
Age	714.0	29.699118	14.526497	0.42	20.1250	28.0000	38.0	
SibSp	891.0	0.523008	1.102743	0.00	0.0000	0.0000	1.0	
Parch	891.0	0.381594	0.806057	0.00	0.0000	0.0000	0.0	
Fare	891.0	32.204208	49.693429	0.00	7.9104	14.4542	31.0	

	max
PassengerId	891.0000
Survived	1.0000
Pclass	3.0000
Age	80.0000
SibSp	8.0000
Parch	6.0000
Fare	512.3292

Show the statistic report of the categorical features of the dataset

	count	unique	top	freq
Name	891	891	Turcin, Mr. Stjepan	1
Sex	891	2	male	577
Ticket	891	681	347082	7
Cabin	204	147	B96 B98	4
Embarked	889	3	S	644

```
[ ]: basicEDA(dfTest, "Titanic Test")
```

Just first five rows

	PassengerId	Pclass	Name	Sex	\
0	892	3	Kelly, Mr. James	male	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	
2	894	2	Myles, Mr. Thomas Francis	male	
3	895	3	Wirz, Mr. Albert	male	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34.5	0	0	330911	7.8292	NaN	Q

1	47.0	1	0	363272	7.0000	NaN	S
2	62.0	0	0	240276	9.6875	NaN	Q
3	27.0	0	0	315154	8.6625	NaN	S
4	22.0	1	1	3101298	12.2875	NaN	S

'The Titanic Test data set consists of 11 different features which for 418 samples.'

Info about the index dtype and columns, non-null values and memory usage.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     418 non-null   int64
1   Pclass          418 non-null   int64
2   Name            418 non-null   object
3   Sex             418 non-null   object
4   Age            332 non-null   float64
5   SibSp           418 non-null   int64
6   Parch          418 non-null   int64
7   Ticket          418 non-null   object
8   Fare           417 non-null   float64
9   Cabin          91 non-null    object
10  Embarked        418 non-null   object
dtypes: float64(2), int64(4), object(5)
memory usage: 36.0+ KB
```

None

Check all of the values in the data frame which holds my data.

	PassengerId	Pclass	Name	Sex \
0	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
1	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
2	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
3	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
4	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
..
413	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
414	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
415	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
416	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>
417	<class 'int'>	<class 'int'>	<class 'str'>	<class 'str'>

	Age	SibSp	Parch	Ticket \
0	<class 'float'>	<class 'int'>	<class 'int'>	<class 'str'>
1	<class 'float'>	<class 'int'>	<class 'int'>	<class 'str'>

```

2    <class 'float'> <class 'int'> <class 'int'> <class 'str'>
3    <class 'float'> <class 'int'> <class 'int'> <class 'str'>
4    <class 'float'> <class 'int'> <class 'int'> <class 'str'>
..    ...
413  <class 'float'> <class 'int'> <class 'int'> <class 'str'>
414  <class 'float'> <class 'int'> <class 'int'> <class 'str'>
415  <class 'float'> <class 'int'> <class 'int'> <class 'str'>
416  <class 'float'> <class 'int'> <class 'int'> <class 'str'>
417  <class 'float'> <class 'int'> <class 'int'> <class 'str'>

```

```

                Fare          Cabin          Embarked
0    <class 'float'> <class 'float'> <class 'str'>
1    <class 'float'> <class 'float'> <class 'str'>
2    <class 'float'> <class 'float'> <class 'str'>
3    <class 'float'> <class 'float'> <class 'str'>
4    <class 'float'> <class 'float'> <class 'str'>
..    ...
413  <class 'float'> <class 'float'> <class 'str'>
414  <class 'float'> <class 'str'> <class 'str'>
415  <class 'float'> <class 'float'> <class 'str'>
416  <class 'float'> <class 'float'> <class 'str'>
417  <class 'float'> <class 'float'> <class 'str'>

```

[418 rows x 11 columns]

Count na values

```

PassengerId    0
Pclass         0
Name           0
Sex            0
Age           86
SibSp          0
Parch         0
Ticket         0
Fare           1
Cabin        327
Embarked       0
dtype: int64

```

"The columns with missing data are: ['Age', 'Fare', 'Cabin']"

```

Percent of missing 'Age' records is 20.574 %
Percent of missing 'Fare' records is 0.239 %
Percent of missing 'Cabin' records is 78.23 %

```

Show the statistic report of the numeric features of the dataset

	count	mean	std	min	25%	50% \
PassengerId	418.0	1100.500000	120.810458	892.00	996.2500	1100.5000
Pclass	418.0	2.265550	0.841838	1.00	1.0000	3.0000
Age	332.0	30.272590	14.181209	0.17	21.0000	27.0000
SibSp	418.0	0.447368	0.896760	0.00	0.0000	0.0000
Parch	418.0	0.392344	0.981429	0.00	0.0000	0.0000
Fare	417.0	35.627188	55.907576	0.00	7.8958	14.4542
	75%	max				
PassengerId	1204.75	1309.0000				
Pclass	3.00	3.0000				
Age	39.00	76.0000				
SibSp	1.00	8.0000				
Parch	0.00	9.0000				
Fare	31.50	512.3292				

Show the statistic report of the categorical features of the dataset

	count	unique	top	freq
Name	418	418	Gracie, Col. Archibald IV	1
Sex	418	2	male	266
Ticket	418	363	PC 17608	5
Cabin	91	76	B57 B59 B63 B66	3
Embarked	418	3	S	270

Some considerations: first of all, we can see how "Survived" is our target: - Survived 0 = no, 1 = yes

As regards the other features we have: - PassengerID: - Pclass: 1 = 1st, 2 = 2nd, 3 = 3rd - Name: self explanatory - Sex: self explanatory - Age: self explanatory - SibSp = nr of sibilings / spouses abroad - Parch = nr of parents / children abroad - Ticket = self explanatory - Fare = passenger fare - Cabin = self explanatory - Embarked = port of embrarkation --> C = Chernourg, Q = Queenstown, S = Southhampton

I think I can do something in order to slim down this dataset.

We can see that the two dataset are similar.

2.1 Feature engineering

So, I think that 'Name', 'Embarked', 'Cabin' and 'Ticket' features can be dropped because I don't believe that be called "Nicholas" or "Augusta" increases the possibility to survive. Same reasoning for the others features. Then, let's start to drop useless features.

```
[ ]: delColumn(dfTrain, "Name")
delColumn(dfTrain, "Ticket")
delColumn(dfTrain, "Cabin")
delColumn(dfTrain, "Embarked")
```

As regards "Fare", let's check out if the fare is related to the better chances to be survive.

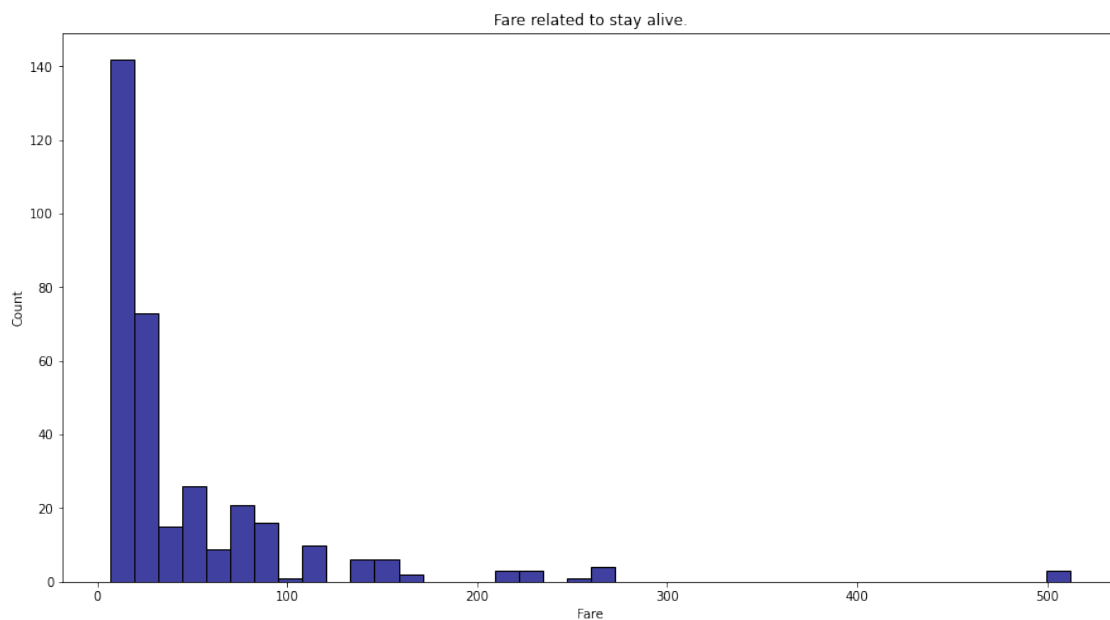
```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns

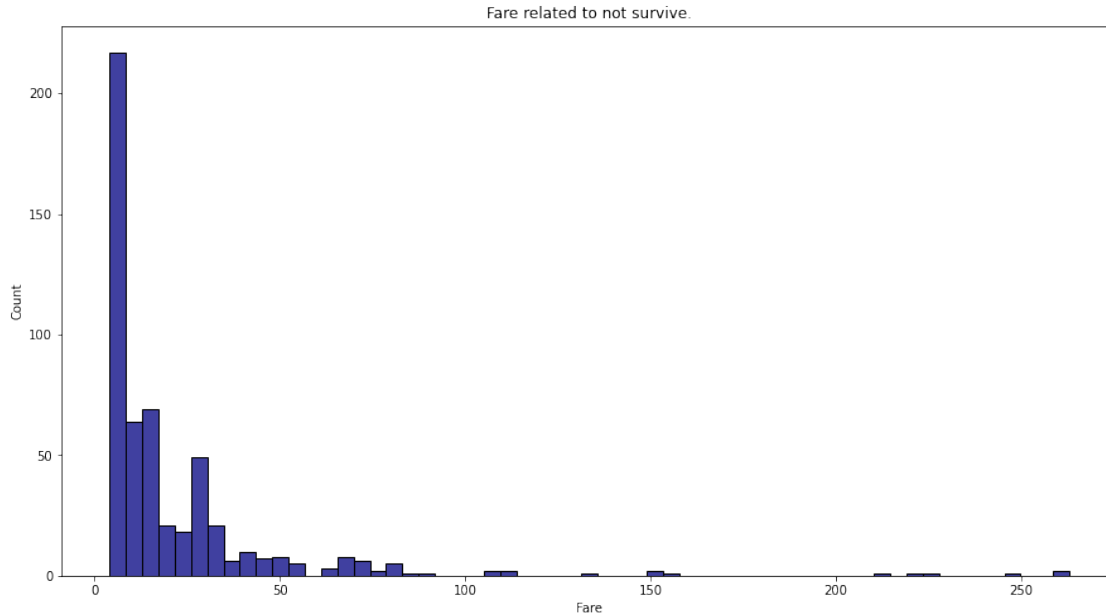
if "Fare" in dfTrain.columns:
    dfTmp = dfTrain[["Fare", "Survived"]]
    dfSurvivedYes = dfTmp[dfTmp.Survived == 1]
    dfSurvivedNo = dfTmp[dfTmp.Survived == 0]

    plt.figure(figsize = (15,8))
    plt.title("Fare related to stay alive.")
    sns.histplot(data = dfSurvivedYes[dfSurvivedYes.Fare > 0],
                  x = 'Fare',
                  color = 'navy'
                  )

    plt.figure(figsize = (15,8))
    plt.title("Fare related to not survive.")
    sns.histplot(data = dfSurvivedNo[dfSurvivedNo.Fare > 0],
                  x = 'Fare',
                  color = 'navy'
                  )

plt.show()
```





So one can be dead or alive, no matter how much he paid as fare: we can also drop this feature.

```
[ ]: delColumn(dfTrain, "Fare")
```

At this point we have few but good features. Remains to resolve the missing 'Age' records (that is 19.865 %). So, let's find out correlations with "Age" feature.

```
[ ]: dfTrain.corr()
```

```
[ ]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch
PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527	-0.001652
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322	0.081629
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081	0.018443
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247	-0.189119
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000	0.414838
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838	1.000000

We have three feature correlate with "Age": Pclass (PCC: -0.369226), SibSp (PCC: -0.308247), Parch (PCC: -0.189119). I'm going to use "IterativeImputer" (which is a multivariate imputer that estimates each feature from all the others) with RandomForestRegressor...

```
[ ]: from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer

from sklearn.ensemble import RandomForestRegressor
import pandas as pd

dftmp = dfTrain.loc[:, ["Age"]]
```

```

imp = IterativeImputer(RandomForestRegressor(),
                        max_iter=10,
                        tol=0.001,
                        random_state=0,
                        sample_posterior=False,
                        verbose=True)
dftmp = pd.DataFrame(imp.fit_transform(dftmp),
                    columns=dftmp.columns)

```

...check if all data are property filled...

```

[ ]: print("\nNumber of rows where 'Age' are null or empty")
     print(dftmp.isnull().sum())

```

```

Number of rows where 'Age' are null or empty
Age      0
dtype: int64

```

...and finally refill the missing Age values.

```

[ ]: delColumn(dfTrain, "Age")
     dfTrain = dfTrain.join(dftmp)

```

Remains to encode the "Sex" feature.

```

[ ]: from sklearn.preprocessing import LabelEncoder

     labelencoder_X = LabelEncoder()

     dfTrain["Sex"] = labelencoder_X.fit_transform(dfTrain["Sex"])

```

So this is our starting dataset.

```

[ ]: dfTrain.head()

```

```

[ ]:

```

	PassengerId	Survived	Pclass	Sex	SibSp	Parch	Age
0	1	0	3	1	1	0	22.0
1	2	1	1	0	1	0	38.0
2	3	1	3	0	0	0	26.0
3	4	1	1	0	1	0	35.0
4	5	0	3	1	0	0	35.0

Same things to test dataset

```

[ ]: delColumn(dfTest, "Name")
     delColumn(dfTest, "Ticket")
     delColumn(dfTest, "Cabin")

```

```

delColumn(dfTest, "Embarked")
delColumn(dfTest, "Fare")

dftmp = dfTest.loc[:, ["Age"]]

imp = IterativeImputer(RandomForestRegressor(),
                        max_iter=10,
                        tol=0.001,
                        random_state=0,
                        sample_posterior=False,
                        verbose=True)
dftmp = pd.DataFrame(imp.fit_transform(dftmp),
                    columns=dftmp.columns)
dfTest.drop("Age", axis=1, inplace=True)
dfTest = dfTest.join(dftmp)

dfTest["Sex"] = labelencoder_X.fit_transform(dfTest["Sex"])

dfTest.head()

```

```

[ ]:
PassengerId  Pclass  Sex  SibSp  Parch  Age
0           892      3    1      0      0  34.5
1           893      3    0      1      0  47.0
2           894      2    1      0      0  62.0
3           895      3    1      0      0  27.0
4           896      3    0      1      1  22.0

```

3 Train and test the model

Once prepared data, we can split data in train and test, as usual.

```

[ ]: from sklearn.model_selection import train_test_split

X = dfTrain.drop("Survived", axis=1)
y = dfTrain["Survived"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4)

```

And prepare the grid search with RandomForestClassifier model, and return

```

[ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, KFold

rndForestParams = {
    "criterion" : ["gini", "entropy"],          # {"gini", "entropy",
    ↪ "log_loss"},

```

```

    "min_samples_leaf" : [1, 5, 10],                # The minimum number of
    ↪ samples required to be at a leaf node.
    "min_samples_split" : [4, 10, 14, 16],          # The minimum number of
    ↪ samples required to split an internal node
    "n_estimators": [150, 300, 700, 1000]           # The number of trees
    ↪ in the forest.
}

rfModel = RandomForestClassifier(
    max_features = "sqrt",                          # The number of
    ↪ features to consider when looking for the best split
    oob_score = True,                               # Whether to use
    ↪ out-of-bag samples to estimate the generalization score.
                                                    # Only available if
    ↪ 'bootstrap = True' (that's default value!)
    random_state = 1,                               # Controls both the
    ↪ randomness of the bootstrapping of the samples used when building trees
    n_jobs = -1                                     # '-1' means using all
    ↪ processors.
)

cv_method = KFold(n_splits = 10, shuffle = True)

gs = GridSearchCV(
    estimator = rfModel,
    param_grid = rndForestParams,
    scoring='accuracy',
    cv = cv_method,
    n_jobs=-1
)

```

Now we can fit the gridsearch object (it will take a while...).

```
[ ]: gs.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=KFold(n_splits=10, random_state=None, shuffle=True),
    estimator=RandomForestClassifier(max_features='sqrt', n_jobs=-1,
                                     oob_score=True, random_state=1),
    n_jobs=-1,
    param_grid={'criterion': ['gini', 'entropy'],
                'min_samples_leaf': [1, 5, 10],
                'min_samples_split': [4, 10, 14, 16],
                'n_estimators': [150, 300, 700, 1000]},
    scoring='accuracy')
```

We can see the parameter setting that gave the best results on the hold out data...

```
[ ]: gs.best_params_
```

```
[ ]: {'criterion': 'entropy',  
      'min_samples_leaf': 1,  
      'min_samples_split': 10,  
      'n_estimators': 150}
```

...and set up a model with the estimator that was chosen by the search.

```
[ ]: RFC_Model = gs.best_estimator_
```

And show what is the average of all cv folds for a single combination of the parameters you specify in the tuned_params.

```
[ ]: gs.best_score_                                     #Mean cross-validated score of u  
      ↪ the best_estimator
```

```
[ ]: 0.8221872816212439
```

Let's predict on train data.

```
[ ]: RFC_Model.predict(X_train)
```

```
[ ]: array([0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0,  
           0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0,  
           0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0,  
           1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1,  
           1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0,  
           0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1,  
           1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,  
           0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,  
           0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,  
           1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,  
           0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,  
           0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,  
           0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0,  
           1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,  
           0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0,  
           0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0,  
           1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0,  
           0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
           0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0,  
           1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1,  
           0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
           1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
           1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0,  
           1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1,  
           0, 0, 0, 0, 0, 0, 0])
```

And make the prediction on test data.

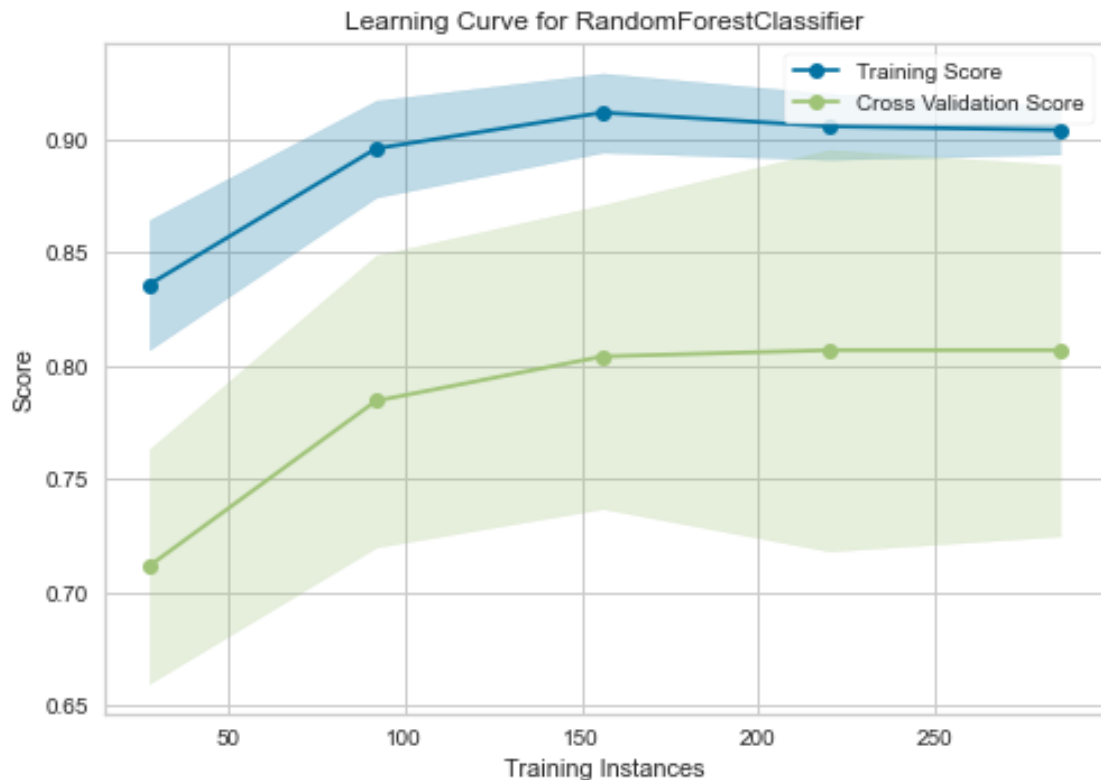
```
[ ]: RFC_Model.score(X_test, y_test) #Return the mean accuracy on  
      ↪ the given test data and labels
```

```
[ ]: 0.8067226890756303
```

3.1 Model Performance Analysis

Displaying the learning curve, we note that we have enough data to try to make a model.

```
[ ]: from yellowbrick.model_selection import learning_curve  
  
      learning_curve(RFC_Model, X_test, y_test, scoring='accuracy')  
      plt.show()
```



```
[ ]: from sklearn.metrics import classification_report  
  
      dfReport = pd.DataFrame(classification_report(y_test, RFC_Model.  
      ↪ predict(X_test), output_dict=True))  
      dfReport
```


```
[ ]:
```

	0	1	accuracy	macro avg	weighted avg
precision	0.808511	0.803279	0.806723	0.805895	0.806415
recall	0.887850	0.685315	0.806723	0.786583	0.806723
f1-score	0.846325	0.739623	0.806723	0.792974	0.803584
support	214.000000	143.000000	0.806723	357.000000	357.000000

```
[ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

predictions = RFC_Model.predict(X_test)
cm = confusion_matrix(y_test, predictions, labels = RFC_Model.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix = cm,
                              display_labels = RFC_Model.classes_)

disp.plot()
plt.grid(visible=None)
plt.show()
```



4 Submission

```
[ ]: predictions = RFC_Model.predict(dfTest)
      predictions
```

```
[ ]: array([0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1,
            1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
            1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1,
            1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1,
            1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
            0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1,
            0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
            0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1,
            1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,
            0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
            1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1,
            0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
            0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0,
            0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,
            0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0,
            1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0,
            0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0,
            1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
            0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0])
```

```
[ ]: PassengerId = dfTest['PassengerId']
      submission = pd.DataFrame({"PassengerId": PassengerId, "Survived": predictions})
      submission.to_csv('submission.csv', index=False)
```